

Partial Fine-Tuning of BioCLIP 2.5 with Taxonomic Auxiliary Heads for Multi-Label Plant Species Identification in Vegetation Plots: ANU at PlantCLEF 2026

Arjun Raj^{1,†}, Manindra de Mel^{1,†} and Razeen Wasif^{1,†}

¹The Australian National University, Canberra, ACT 2601, Australia

Abstract

This paper describes the system submitted by the Australian National University team to the PlantCLEF 2026 challenge, which tasks participants with predicting all plant species present in high-resolution vegetation quadrat images. The core difficulty lies in the domain shift between the training data, comprising individually photographed single-plant specimens, and the test data, consisting of dense multi-species vegetation plots. Our final submission centres on a partial fine-tune of BioCLIP 2.5 (ViT-H/14) in which only the last four transformer blocks together with the final layer norm and projection are unfrozen, while the lower twenty-eight blocks are kept frozen to preserve the Tree-of-Life prior. The encoder is trained on the enlarged PlantCLEF 2024 single-plant manifest with auxiliary cross-entropy heads at the genus, family, and where available order and class taxonomic levels, supplying a structural regulariser without adding inference-time cost. Inference decomposes each quadrat into a 4×4 grid of 448-pixel tiles with no overlap, forwards each tile through the encoder at both 224 and 336 pixels, and aggregates the tile distributions by per-species softmax mean across the sixteen tiles and across the two resolutions; class-prior logit adjustment ($\tau=0.25$) is then applied and an adaptive probability threshold ($T=0.03$, $k_{\min}=2$, $k_{\max}=10$) selects the emitted species set. Beyond this anchor we report three architecture-orthogonal pivots that we submitted but that did not improve over the visual posterior: a triple-backbone classifier retrained on frozen BioCLIP 2.5, DINOv3, and ConvNeXt-V2-L features; a genus and family co-occurrence rerank; and a seasonal phenology prior derived from GBIF month-of-observation histograms. Across all four submitted configurations, our best public Macro F1 is **0.418265** and our best private Macro F1 is **0.40283**.

Keywords

plant identification, multi-label classification, partial fine-tuning, BioCLIP, taxonomic auxiliary loss, vegetation plot analysis, PlantCLEF 2026

1. Introduction

Monitoring plant biodiversity at scale is essential for understanding ecosystem health, tracking the effects of climate change, and informing conservation policy. Traditionally, vegetation surveys require expert botanists to visit field sites, place standardised quadrats (typically 50×50 cm sampling frames), and manually identify every plant species visible within the frame. This process is labour-intensive, difficult to scale, and constrained by the availability of trained taxonomists [1]. The PlantCLEF 2026 challenge [2] seeks to automate this task by framing it as a multi-label classification problem: given a high-resolution photograph of a vegetation quadrat, predict the set of all plant species present. The challenge draws from a taxonomy of approximately 7,800 candidate species, and submissions are evaluated using a macro-averaged F1 score computed per sample and then aggregated across transects. A central difficulty of this challenge is the *domain shift* between the available training data and the target test data. The training set consists of over one million single-plant images from the PlantCLEF 2024 corpus, each depicting an individual specimen with a single species label (Figure 1). In contrast, the test images are dense, complex vegetation scenes where multiple species overlap, occlude one another, and appear at varying scales and growth stages (Figure 2). Models trained exclusively on isolated specimens tend to perform poorly when confronted with these multi-species scenes, as they have never encountered the visual complexity of real-world vegetation during training.

CLEF 2026 – Conference and Labs of the Evaluation Forum, September 21–24, 2026, Jena, Germany

[†]These authors contributed equally.

✉ u7526852@anu.edu.au (A. Raj); u7543337@anu.edu.au (M. d. Mel); u7619822@anu.edu.au (R. Wasif)



© 2026 Copyright for this paper by its authors. Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0).



Figure 1: Six representative single-plant training images from the PlantCLEF 2024 corpus. Each image carries a single species label and was contributed through the Pl@ntNet citizen-science pipeline. Shooting style is intentionally varied (close-up macro, half-body, full-plant in habitat) and lighting / background are uncontrolled.



Figure 2: Six representative quadrats from the PlantCLEF 2026 test set, illustrating the breadth of within-set domain variation. Filenames (left to right): CBN-Pd1C-A1-20130807, CBN-Pla-A1-20130808, GUARDEN-CBNMed-10-1-16-34-20240429, LISAH-JAS-0-1-20220505, OPTMix-0108-P1-312-20231006, RNNB-1-1-20230512. Each filename encodes the transect site and observation date. The quadrats range from dense mid-summer canopies (CBN-Pd1C, RNNB) to dry shrub-mat communities (LISAH, GUARDEN) and litter-covered shoulder-season frames (OPTMix). Shooting angle, sampling-frame style (wood, tape, painted line), and illumination all vary across the corpus, compounding the single-plant $\rightarrow$ multi-species domain shift seen in Figure 1.

A second major challenge is the *long-tail distribution* of the training data. While common species may have thousands of training images, rare species—which are often the most ecologically important—can have fewer than ten. Standard cross-entropy training causes models to overwhelmingly predict common classes, yielding high micro-accuracy but poor macro-F1. In this paper, we present the system submitted by the Australian National University (ANU) team. Our approach centres on three key design decisions:

1. **Partial-unfreeze BioCLIP 2.5 with taxonomic regularisation.** BioCLIP [3] is a contrastive vision-language model whose pretraining on TreeOfLife-10M aligns its feature space with the Linnaean taxonomy. Rather than fine-tuning the full backbone—which we show collapses this prior—we unfreeze only the last four transformer blocks of BioCLIP 2.5 ViT-H/14 and attach auxiliary cross-entropy heads at the genus, family, order, and class levels. The auxiliary losses are weighted by their entropy (coarser levels receive smaller weights) so that they primarily regularise the encoder rather than drive it.
2. **Long-tail data rebalancing.** The PlantCLEF 2024 single-plant corpus is heavily skewed: a small number of species exceed one thousand training images while hundreds carry fewer than ten. We attack this from both ends, by augmenting the manifest with extra observations for under-represented species and by applying a per-species image cap that flattens the head of the distribution. We find that on this benchmark, more training data and a flatter species distribution dominate gains from additional unfrozen capacity.
3. **Tile-level inference with adaptive selection.** At inference we partition each quadrat into a 4×4 grid of 448-pixel tiles with no overlap, forward each tile through the encoder at both 224 and 336 pixels, and aggregate the per-tile softmax distributions by per-species mean across both the sixteen tiles and the two resolutions. Class-prior logit adjustment ($\tau=0.25$) is then applied.

Final selection is adaptive: we emit every species whose post-adjusted probability exceeds a fixed threshold ($T=0.03$), with a floor of two and a ceiling of ten species per quadrat. The 224 + 336 resolution ensemble lifts the public score by +0.007 over the single-resolution baseline; the adaptive threshold outperforms fixed top- k selection by between 0.005 and 0.008 Macro F1.

The remainder of this paper is organised as follows. Section 2 reviews related work on plant identification, foundation model adaptation, and multi-label classification. Section 3 describes our system architecture and training procedure in detail. Section 4 presents our experimental setup and results. Section 7 concludes with a discussion of limitations and directions for future work.

2. Related Work

2.1. Automated Plant Identification

Large-scale plant identification has been driven by the LifeCLEF series of evaluation campaigns [1], which have progressively increased in scale and difficulty. Early systems relied on hand-crafted features such as leaf shape descriptors and colour histograms [4]. The advent of deep convolutional networks brought substantial improvements, with Inception and ResNet architectures achieving competitive results on single-specimen classification [5]. More recent work has shifted toward Vision Transformers (ViTs), which excel at capturing long-range spatial dependencies within images [6]. The PlantCLEF 2024 and 2025 editions introduced the vegetation plot setting, highlighting the domain gap between specimen-level training data and plot-level test data [7]. Top-performing systems in those editions employed strategies such as tiled inference (SAHI-style slicing), test-time augmentation, and ensemble methods to bridge this gap.

2.2. Foundation Models for Biodiversity

BioCLIP [3] is a contrastive vision-language model trained on the BIOSCAN and TreeOfLife-10M datasets, which encompass millions of organism images annotated with hierarchical taxonomic labels. Unlike general-purpose CLIP models, BioCLIP aligns visual features with the Linnaean taxonomy, making its feature space inherently suited for biological classification tasks. Studies have shown that BioCLIP features outperform ImageNet-pretrained features for species identification, particularly for morphologically similar species [3]. DINOv3 [8] is a self-supervised ViT trained with a student-teacher distillation objective on a large curated dataset. Its features are highly transferable across domains and exhibit strong spatial correspondence, making them effective for fine-grained recognition tasks. We use DINOv3 features as one of three frozen feature extractors in our triple-backbone pivot (Section 4). ConvNeXt-V2 [9] modernises the convolutional neural network paradigm by incorporating design elements from Transformers (e.g., larger kernel sizes, LayerNorm, GELU activations) alongside a Fully Convolutional Masked Autoencoder (FCMAE) pretraining strategy. ConvNeXt-V2 retains the inductive biases of CNNs—translation equivariance and locality—while achieving performance competitive with ViTs.

2.3. Partial Fine-Tuning of Large Pretrained Models

Fine-tuning all parameters of large pretrained models is often impractical: it requires substantial memory, risks catastrophic forgetting of pretraining knowledge, and tends to overfit when the downstream training set is noisy or long-tailed. A widely used alternative is to freeze the bulk of the backbone and adapt only the top few transformer blocks, which preserves the lower-layer features that capture domain-general structure while allowing the upper layers to specialise. For taxonomically aligned foundation models such as BioCLIP, this trade-off is particularly favourable, since the pretraining prior itself is what makes the model valuable on biological tasks.

2.4. Multi-Label Classification and Long-Tail Learning

Multi-label classification requires predicting a set of labels per input, typically using sigmoid activations and binary cross-entropy. Asymmetric Loss (ASL) [10] extends focal loss [11] to the multi-label setting by applying separate focusing parameters γ_+ and γ_- for positive and negative samples. By setting $\gamma_- > \gamma_+$, ASL aggressively down-weights easy negatives, which dominate in settings with thousands of classes. Logit adjustment [12] provides a complementary approach by shifting logits based on class prior probabilities, effectively requiring the model to produce larger margins for common classes and smaller margins for rare ones.

3. Methodology

3.1. Project Pipeline Overview

Our submission is the product of eighteen numbered experiments and roughly 240 leaderboard submissions converging on a single supervised fine tune of BioCLIP 2.5 ViT-H/14. Starting from a frozen prototype baseline (experiment 006) at 0.330 Macro F1, we briefly explored full fine tuning (experiment 009), which collapsed the Tree of Life prior and dropped the model to 0.208. We then narrowed the trainable surface to the last four transformer blocks together with the final layer norm and projection layer, retaining the lower twenty eight blocks in their pretrained state. This partial unfreeze recipe (experiment 010), augmented with auxiliary classification heads at genus, family, order, and class taxonomic levels, became our 0.38333 anchor. The track culminated in experiment i002, which ported the same partial unfreeze recipe onto a substantially enlarged training manifold (2.65 million images, drawn from the redownloaded PlantCLEF 2024 corpus with pre-filled genus and family labels) and trained for ten epochs instead of five. The best checkpoint (last-four-blocks at epoch ten) was deployed for inference under two recipes that are both reported in Section 4: a single-resolution submission at $T=0.05$ scoring 0.41165, and an ensemble of inference at 224 and 336 pixels combined with logit adjustment ($\tau=0.25$) and $T=0.03$ scoring **0.418265**. The latter is our headline submission; its private Macro F1 is **0.40283**. A follow-up (experiment i003) added 161,686 extra observations for the 2,022 species with fewer than one hundred training images and capped the resulting manifest at 500 images per species, but the resulting checkpoints did not surpass the uncapped i002 best. Beyond the anchor, we submitted three architecture-orthogonal pivots that aimed to add information sources the visual encoder could not access from a single tile: a triple-backbone classifier retrained on frozen BioCLIP 2.5, DINOv3, and ConvNeXt-V2-L features; a genus and family co-occurrence rerank derived from the BioCLIP 2.5 posterior itself; and a seasonal phenology prior built from GBIF month-of-observation histograms keyed to each quadrat’s observation date. All three are reported in Section 4; none improved over the visual anchor. All submissions share the same inference scaffold, illustrated in Figure 3: SAHI-style four by four tiling at tile size 448 with no overlap, softmax mean aggregation across the sixteen tiles, and adaptive selection that emits between two and ten species per quadrat at a probability threshold T . The headline anchor adds two further ingredients: an ensemble of inference at 224 and 336 pixel resolutions, and logit adjustment ($\tau=0.25$) before threshold selection at $T=0.03$. The exact T used by each pivot is reported alongside that pivot.

3.2. Partial Unfreeze BioCLIP 2.5 with Taxonomic Auxiliary Heads

The 010 fine tune treats BioCLIP 2.5 as a taxonomically aligned feature extractor whose lower twenty eight transformer blocks already encode the inductive bias that Tree of Life pretraining provides. Only the last four blocks, the final layer norm, and the projection head are unfrozen. This narrow surface is sufficient for adapting the encoder to the PlantCLEF 2024 single plant distribution without destroying the broader biological prior, and it is also resilient to optimiser noise on a long tailed training set. The head consists of a shared multi layer perceptron mapping the 1,024 dimensional CLS embedding to a hidden representation through layer normalisation, a single linear projection, GELU activation, and dropout. From this representation we attach linear classification heads, one per taxonomic level.




Figure 3: Inference pipeline for the headline anchor submission (solid path; public Macro F1 0.418265). Each quadrat is partitioned into a 4x4 grid of 448-pixel non-overlapping tiles and forwarded through the partially fine-tuned BioCLIP 2.5 ViT-H/14 (lower 28 transformer blocks frozen; last 4 plus LN and projection unfrozen) at two inference resolutions, 224 and 336 pixels. Per-tile softmax probabilities are averaged across the 16 tiles and across the two resolutions, then class-prior logit adjustment ($\tau=0.25$) is applied before an adaptive sigmoid threshold ($T=0.03$, $k_{\min}=2$, $k_{\max}=10$) selects the output species set. Three architecture-orthogonal pivots (dashed; cRT triple-backbone classifier retraining, genus co-occurrence rerank, seasonal phenology prior) indicate where alternative configurations insert into the anchor; all three scored below the anchor on the leaderboard and are reported in Section 4.

Experiment 010 uses the original PlantCLEF 2024 metadata merged with a GBIF lookup and exposes five heads with output dimensions 7,806 for species, 1,446 for genus, 181 for family, 61 for order, and 6 for class. Experiment i002 uses the redownloaded i001 manifest, which carries species, genus, and family columns but no order or class columns; the model therefore drops the two coarsest heads. The training objective is a weighted joint cross entropy

$$\mathcal{L} = \mathcal{L}_{\text{species}} + \sum_{l \in L} w_l \mathcal{L}_l, \quad (1)$$

where the species term carries an implicit weight of one, $L = \{g, f, o, c\}$ for 010 and $L = \{g, f\}$ for i002, and the auxiliary weights are $w_g=0.30$, $w_f=0.15$, $w_o=0.05$, $w_c=0.02$. The smaller weights on coarser levels reflect their lower entropy: when present, the order and class heads converge quickly and act primarily as regularisation, biasing the encoder toward representations whose taxonomic structure is internally consistent. Missing taxonomy labels (NaN in the CSV) are encoded as -1 and masked out of the auxiliary loss. Both runs train in bfloat16 across two RTX 5090 GPUs, with a backbone learning rate of 1×10^{-6} , a head learning rate of 1×10^{-4} , one epoch of linear warmup followed by cosine decay, weight decay 1×10^{-4} , and label smoothing 0.1. Experiment 010 trains for five epochs at batch size 64 per replica with gradient accumulation of four (effective batch 512). Experiment i002 trains for ten epochs on the larger manifest at batch size 128 per replica with gradient accumulation of four (effective batch 1024); the longer schedule is needed because the training set is roughly $1.7 \times$ larger. An ablation reported in Section 4 fixes the 010 recipe and varies only the number of unfrozen blocks; performance peaks sharply at four blocks and degrades by 0.009 to 0.014 Macro F1 at three and five blocks respectively.

3.3. Long Tail Rebalancing: More Data vs. Per-Species Capping

The PlantCLEF 2024 single plant corpus is heavily skewed. Of the 7,806 species, the most represented twelve carry more than one thousand training images while one hundred forty five species carry a single image and four hundred forty five carry between three and five. Vanilla shuffled training thus exposes the network to roughly two orders of magnitude more head class examples than tail class examples per epoch. Experiments i002 and i003 attack this imbalance along two distinct axes, the second building directly on the first. **Experiment i002 – more data, no cap.** Experiment i002 ports the 010 partial unfreeze recipe onto the redownloaded i001 manifest, which carries 2,653,781 rows across 7,806 species (against 1,408,033 rows in 010) with pre-filled genus and family labels for every row. No per species cap is applied; the rebalancing here is implicit, coming from the broader observation pipeline rather than the tail. The stratified 10% validation split yields 2,391,457 training rows and 262,324 validation rows. The downstream effect on validation accuracy is substantial: top 5 accuracy on the held out split rises from 0.9214 (010 head only) and 0.9323 (010 last four blocks) to 0.9505 (i002 head only), 0.9583 (i002 last four blocks), and 0.9615 (i002 last eight blocks). **Experiment i003 – capped manifest with tail augmentation.** Experiment i003 takes the i002 manifest and explicitly attacks both ends of the long tail. We augment it with an additional 161,686 validated images sourced for 2,022 species with fewer than one hundred training examples, decoded through a full Python Imaging Library verification pass to remove truncated, non-image, or oversized files. The augmented manifest is then deduplicated by image path and capped at five hundred images per species through a seeded random subsample, applied once offline at manifest construction time rather than at every epoch. This trims the heaviest bins (the twelve species above one thousand images and the 2,263 species in the 501 to 1,000 image bin both collapse into the ≤ 500 bin) and produces a flatter distribution. **Result.** The downstream effect on the public leaderboard is reported in Section 4, and contains an empirical surprise: capping at five hundred images per species *hurt* relative to letting the model see the full uncapped manifold. The i003 best is 0.40041 on the leaderboard versus 0.418265 for the corresponding uncapped i002 checkpoint. We interpret this as direct evidence that even on a long tail, the marginal value of additional head class examples is positive on this benchmark provided their label quality is high; flattening the head distribution discards information rather than rebalancing it.




Figure 4: Species count per training-image bin, before and after the i003 per-species cap at 500. The cap collapses the two heaviest bins ($\leq 1k$ and $> 1k$) into the ≤ 500 bin and lifts middle bins, but does not move species out of the long tail (left bins). Section 3.3 reports the corresponding leaderboard outcome: this flatter distribution *hurt* the public Macro F1.

3.4. Adaptive Tile Inference

At inference time each quadrat is resized to a 1792×1792 pixel canvas and partitioned into a four by four grid of non-overlapping 448 pixel tiles. The fine-tuned encoder is then applied to each tile under our headline recipe at two resolutions: tiles downsampled to the encoder’s native 224 pixels, and tiles

downsampled to 336 pixels (the BioCLIP 2.5 ViT-H/14 backbone tolerates the $1.5\times$ stretch via bicubic positional-embedding interpolation). The per-tile species logits are mapped to probabilities via softmax, the sixteen tile distributions are aggregated by per-species mean, and the resulting 224 and 336 posteriors are averaged. Softmax-mean aggregation was selected after an ablation against per-species max, mean of top three, and noisy-or; mean was the most numerically stable on the long tail and dominates max aggregation by roughly 0.04 Macro F1 on the i002 last-four-blocks checkpoint. We replace the standard top k selection with an adaptive thresholded rule combined with class-prior logit adjustment. After aggregation we apply logit adjustment $\ell'_s = \ell_s - \tau \log \tilde{\pi}_s$ with $\tau=0.25$, where $\tilde{\pi}_s$ is the Laplace-smoothed training prior for species s , and emit every species whose post-adjusted probability exceeds $T=0.03$, subject to a floor of $k_{\min}=2$ and a ceiling of $k_{\max}=10$. A sweep of $T \in \{0.02, 0.03, 0.05\}$ together with several gap and relative thresholds was performed on each checkpoint; the operating point for the headline 224+336 ensemble is $T=0.03$. For ablation context, the same checkpoint at single-resolution 224 pixels *without* logit adjustment scores 0.41165 at $T=0.05$; the 224+336+LA ensemble lifts this to 0.418265, a gain of +0.007. Adaptive selection contributes a further 0.005 to 0.008 Macro F1 over the corresponding fixed top-three submission on the same checkpoint.

3.5. Seasonal Phenology Prior

Each test quadrat carries an observation date in its identifier (e.g. CBN-Pd1C-A1-20130807 $\rightarrow$ 7 Aug 2013); the visual backbone never sees this. Date is therefore the cleanest source of information that is *information-theoretically orthogonal* to the image-based posterior. We formalise a phenological prior as a circular probability density over day-of-year, estimated nonparametrically from GBIF month-of-observation histograms and combined with the visual posterior multiplicatively in log space.

Formulation. For a species s observed on day-of-year $d \in \{1, \dots, 365\}$, let $P(s \mid d)$ denote the species' observability density. We define the combined posterior as

$$\log p_{\text{final}}(s \mid I, d) = \log p_{\text{vis}}(s \mid I) + \beta \log P(s \mid d), \quad (2)$$

where p_{vis} is the BioCLIP 2.5 anchor posterior and β is a mixing weight: $\beta=0$ recovers the unmodified visual posterior; large β collapses the prediction toward the seasonal prior alone.

Pdf estimation. We estimate the per-species pdf nonparametrically from GBIF month-of-observation histograms. The histograms are obtained via a single faceted query per species:

```
GET /v1/occurrence/search?taxonKey=K
    &facet=month&facetLimit=12&limit=0
```

which returns a 12-bin month histogram in one request. We then build a 365-day pdf via circular Gaussian smoothing with $\sigma = 18$ d centred on each month's mid-point, and mix in a small uniform mass to prevent hard zeros in tail months:

$$P(s \mid d) = (1-\epsilon) \frac{1}{Z_s} \sum_{m=1}^{12} c_{s,m} \mathcal{N}_{\text{circ}}(d \mid \mu_m, \sigma) + \frac{\epsilon}{365}, \quad (3)$$

with $\epsilon=0.05$ and $c_{s,m}$ the GBIF count for species s in month m .

Coverage. We fetched histograms for all 7,796 species with valid GBIF identifiers (10 of 7,806 lack a GBIF mapping). Two passes were required: a first pass at 12 worker threads (3 req/s, 21% rate-limit drops), then a retry pass at 4 worker threads (GBIF aggressively throttled a saturated client). Final coverage: **6,957 / 7,806 species (89.2%)**; the remaining 849 species fall back to a uniform DOY pdf and contribute no phenological signal.




Figure 5: Inference pipeline for the seasonal phenology submission (Pivot 3, public Macro F1 0.41346). Each quadrat is processed at two grid scales (1.0 and 0.8); tiles are screened by an Excess Green (ExG) vegetation filter; the surviving tiles are forwarded through the partially fine-tuned BioCLIP 2.5 ViT-H/14 at two inference resolutions (224 and 336 pixels). Per-tile probabilities are pooled by entropy-weighted Bayesian aggregation, $w_t \propto \exp(-H_t) \cdot \text{ExG}_t$, and the two resolutions are then averaged. A Boltzmann-weighted seasonal phenology prior over the GBIF day-of-year histogram is added in log space at $\beta=1.0$, followed by class-prior logit adjustment ($\tau=0.25$) and adaptive sigmoid threshold selection ($T=0.03$, $k_{\min}=2$, $k_{\max}=10$). The pipeline differs from the anchor (Figure 3) in its multi-scale tiling, ExG screening, and entropy-weighted aggregation; the phenology log-prior is the additional term unique to this submission.

Submission pipeline. Beyond the phenology log-prior itself, the seasonal phenology submission departs from the anchor inference scaffold (Figure 3) in three ways, illustrated architecturally in Figure 5 and visually in Figure 6. First, an Excess Green index $\text{ExG}_t = 2G_t - R_t - B_t$ is computed per tile and tiles with fewer than 15% green pixels are dropped before the encoder forward pass, removing tiles that contain no vegetation. Second, the surviving tiles are aggregated with a per-tile Bayesian weight $w_t \propto \exp(-H_t) \cdot \text{ExG}_t$, where H_t is the predictive entropy of tile t , yielding a softmax-weighted mean rather than the anchor’s unweighted mean. Third, the input quadrat is processed at two grid scales (1.0 and 0.8) in addition to the two inference resolutions (224 and 336 pixels), with all per-tile distributions pooled into a single posterior before the phenology prior is added. Logit adjustment ($\tau=0.25$) and adaptive sigmoid threshold selection ($T=0.03$) then follow as in the anchor.

Empirical result. On the PlantCLEF 2026 leaderboard, the phenology prior at $\beta=1.0$, $T=0.03$ produced a Macro-F1 of **0.41346** versus the unaugmented BioCLIP 2.5 anchor of 0.418265 ($\Delta =$

−0.0048). The detailed sweep and discussion are reported in §4.




Figure 6: Data flowing through the seasonal phenology submission, stage by stage. *Multi-scale tiling:* the input quadrat is partitioned into a 4×4 grid at the native scale (solid lines) and re-partitioned at scale 0.8 (dashed inner rectangle), giving two tile sets. *ExG filter:* each tile's Excess Green index $\text{ExG}_t = 2G_t - R_t - B_t$ is computed and tiles with fewer than 15% green pixels are dropped (greyed out with a cross). *Entropy weights:* the surviving tiles are weighted by $w_t \propto \exp(-H_t) \cdot \text{ExG}_t$, where bar height encodes the weight; confident, vegetation-rich tiles dominate the softmax-weighted mean. *Phenology curve:* the per-species observability density $P(s | d)$, smoothed from the GBIF month-of-observation histogram, is multiplied into the posterior in log space at $\beta=1.0$, where the dashed red line marks the quadrat's observation date d . *Species selection:* after class-prior logit adjustment, species whose probability exceeds $T=0.03$ are emitted; the bottom bar (greyed) falls below threshold and is rejected.

3.6. Considered Approaches Not Implemented

During the project we scoped a number of additional directions that did not make it into the submitted system, either because the engineering cost outweighed the expected gain on this benchmark or because preliminary work indicated negative or noise-floor deltas. We list them briefly here for completeness, without method-level detail. Gated multi-backbone fusion with a learned per-image gating network over BioCLIP 2.5, DINOv3, and ConvNeXt-V2-L was prototyped but never completed end-to-end; an offline-cached high-resolution knowledge distillation pipeline with two SWA-averaged frozen teachers was specified but not trained; a graph neural network head over a blended ecological and phylogenetic adjacency matrix was designed and partially implemented; Markov random field reasoning over quadrat tiles via loopy belief propagation was scoped; retrieval-augmented classification using per-species feature prototypes was scoped for the extreme tail; test-time training via self-supervised rotation prediction was scoped; conformal prediction for set-valued species output was scoped; and Frank-Wolfe optimisation with an island-biogeography count prior was scoped as an alternative to the threshold-based selection rule. None of these are evaluated in this paper.

4. Experiments

4.1. Anchor Configuration

Our anchor configuration is the BioCLIP 2.5 last four blocks tile-inference submission described in Section 3.4: 4×4 tiling at tile size 448 with no overlap, an ensemble of inference at 224 and 336 pixel resolutions, softmax mean aggregation per tile, logit adjustment at $\tau=0.25$, and probability threshold selection at $T=0.03$ with a floor of two and a ceiling of ten species per quadrat. This configuration scores **0.418265** public Macro F1 on the PlantCLEF 2026 leaderboard (i002 four-blocks 224+336 ensemble; private Macro F1 = 0.40283). For comparison, the same checkpoint at single-resolution 224 pixels without logit adjustment (`last_blocks_infer_softmax_mean / softmax_mean_probT0.05`) scores 0.41165 public / 0.38132 private — the 224+336+LA ensemble lifts the public score by +0.007 and the private score by +0.022.

4.2. Unfrozen-Blocks Ablation

The number of unfrozen transformer blocks is a sharp, not a plateau, design variable. We ran an ablation on the 010 recipe (the original 1.4M image manifest, five epochs, identical optimiser and head configuration), varying only the number of unfrozen blocks. Table 1 reports the resulting public Macro F1. The head-only baseline corresponds to experiment 006 and the all-blocks (full fine-tune) row to experiment 009.

Unfrozen blocks	Public Macro F1
0 (head only)	0.330
3 blocks	0.374
4 blocks	0.38333
5 blocks	0.369
All 32 blocks (full FT)	0.208

Table 1

Unfrozen-blocks ablation on the 010 recipe. Performance peaks sharply at four blocks; one block too few costs 0.009 Macro F1, one block too many costs 0.014. Full fine-tuning collapses the Tree-of-Life prior and is catastrophic.




Figure 7: Unfrozen-blocks sweep visualised with both metrics. Held-out validation top-5 (dashed, right axis) rises monotonically with capacity from $n=0$ to $n=5$, but Kaggle Macro F1 (solid, left axis) peaks sharply at $n=4$ and collapses at full fine-tune ($n=32$). The disagreement between the two curves is the load-bearing finding of §5: validation accuracy on the single-plant distribution does not predict multi-species quadrat Macro F1.

4.3. Within-Backbone Saturation

Before pivoting to architecture-orthogonal evidence sources, we measured whether further within-backbone diversity could break the i002 anchor by comparing the four-blocks i002 checkpoint against the eight-blocks i002 checkpoint and against epoch-5 versions of both. Table 2 reports the pairwise top-one agreement and submission Jaccard on 2,105 test quadrats. The 0.923 Jaccard between the four

Pair	Top-1	Jaccard
blk_4 vs. blk_8	0.901	—
blk_4 vs. blk_4 (epoch 5)	0.807	0.664
4-way intra-BC2.5 ens. vs. anchor	—	0.923

Table 2

Within-backbone diagnostic gates. Top-1 agreement on a 7,806-way classifier and submission Jaccard at the operating threshold.

way intra-BioCLIP 2.5 ensemble and our anchor submission demonstrates that all within-backbone

variants are statistically equivalent on the test set: shuffling the unfrozen depth or the training-epoch index does not move the prediction distribution. This motivates the pivot to fundamentally different evidence sources.

4.4. Pivot 1: Triple-Backbone Classifier Retraining (cRT)

We trained two MLP classifier heads on a frozen 3328-d concatenation of BioCLIP 2.5 ViT-B, DINOv3 ViT-L, and ConvNeXt-V2-L features (head A: standard CE; head B: class-balanced random sampling). Best validation accuracy: 63.95%. We applied the trained heads at inference using the same 4×4 tile / softmax mean aggregation pipeline as the BioCLIP 2.5 anchor. **Pre-submission diagnostic gate.** Before submitting, we evaluated structural orthogonality against the anchor using top-one agreement and submission Jaccard, a cheap diagnostic that does not consume leaderboard submissions. Top-one agreement between the cRT head and the BioCLIP 2.5 anchor was only 0.094 (vs. 0.807 for two BioCLIP 2.5 checkpoints of the same family), and submission Jaccard at $T=0.05$ was 0.078. We view this agreement-Jaccard screen as a reusable methodological gate for working-note teams: it lets a candidate ensemble be characterised as either redundant (Jaccard near 1) or structurally divergent (Jaccard near 0) before any leaderboard budget is spent. **Submission.** Mixing the head A / head B mean of cRT with the BioCLIP 2.5 anchor at $w=0.3$, $T=0.03$, $\tau=0.25$ scored **0.38195** Macro F1, $\Delta = -0.036$ versus the 0.418265 anchor. **Reading.** cRT is structurally orthogonal but *empirically wrong* on the test distribution: a peakier posterior (mean $k=1.33$ standalone) means high-confidence wrong picks dominate the ensemble. The 9% top one agreement was diversity going *against* signal rather than complementing it.

4.5. Pivot 2: Genus / Family Co-occurrence Reranking

For each quadrat, we computed aggregated genus and family supports

$$P_{\text{gen}}(g) = \sum_{s \in g} p_{\text{BC}}(s), \quad (4)$$

$$P_{\text{fam}}(f) = \sum_{s \in f} p_{\text{BC}}(s), \quad (5)$$

where $p_{\text{BC}}(s)$ is the BioCLIP 2.5 anchor posterior, and applied a “siblings stick together” log-prior:

$$\log p_{\text{new}}(s) = \log p_{\text{BC}}(s) + \alpha_g \log P_{\text{gen}}(g(s)) + \alpha_f \log P_{\text{fam}}(f(s)). \quad (6)$$

We swept $\alpha_g \in \{0, 0.1, 0.2, 0.3, 0.5\}$, $\alpha_f \in \{0, 0.05, 0.1, 0.2\}$, $T \in \{0.025, 0.03, 0.035\}$ and submitted the *isovolumetric* candidate $\alpha_g=0.1$, $\alpha_f=0.0$, $T=0.035$, $\tau=0.25$ (mean $k=2.89$, matched to the BioCLIP 2.5 anchor volume; submission Jaccard with the anchor = 0.887). Score: **0.41246**, $\Delta = -0.006$. **Reading.** The rerank swapped 11% of selections to genus-supported siblings at constant volume; the score moved by only -0.006 , comfortably inside the run-to-run noise floor. Priors derived from the existing logit structure (genus support is a reweighted sum of probabilities the model already produced) carry essentially no new information – the same evidence repackaged.

4.6. Pivot 3: Seasonal Phenology Prior

The full methodology is described in §3.5 and the pipeline is illustrated in Figure 5. We swept $\beta \in \{0.25, 0.5, 1.0, 1.5, 2.0\}$ at three probability thresholds. The natural Bayesian operating point ($\beta=1.0$, no temperature scaling, $T=0.03$ and $\tau=0.25$ matching the anchor’s logit-adjustment and threshold values) flips 12.8% of top one picks at mean $k=3.24$. Score: **0.41346**, $\Delta = -0.005$. The DOY distribution of the test set is heavily summer skewed (Jul = 517, Aug = 668, Apr–Jun = 607 of 2,105 quadrats), confirming that phenology should be most informative on summer flowering taxa. Figure 8 visualises




Figure 8: Calendar wheel of the PlantCLEF 2026 test set by month of observation (2,105 quadrats total). Radial bar length is proportional to the count for that month. The July (517) and August (668) wedges are reported exactly in the data; the April–June total of 607 is split uniformly across the three months for the figure, as is the residual 313 across the remaining seven months. The visually obvious summer over-representation motivates the seasonal phenology pivot (§4.6): a date-conditioned prior could only ever help in the densely-sampled months, and conversely should not be expected to lift Macro F1 on the sparse winter tail.

this concentration as a calendar wheel: roughly 85% of test quadrats fall between April and August, leaving the remaining seven months almost empty.

Reading. Three concurrent factors plausibly explain the near zero delta: (i) GBIF observation effort is itself summer biased, diluting the discriminative season signal; (ii) the BioCLIP 2.5 backbone trained on dated images already encodes seasonal phenotype implicitly via leaf / flower / fruit state; (iii) 11% of species fall back to a uniform pdf precisely on the long tail where the visual model is most uncertain — so phenology cannot break ties where it would be most useful.

4.7. Inference Hyperparameter Ablations

The anchor and unfrozen-blocks (§4.2) results commit to specific values for several inference-side knobs: logit-adjustment temperature $\tau=0.25$, probability threshold $T=0.03$, softmax-mean tile aggregation, and an adaptive probability-threshold selection rule with $k_{\min}=2$, $k_{\max}=10$. This subsection consolidates the leaderboard sweeps that those values fell out of. All numbers below are public Macro F1; each table or paragraph names the underlying checkpoint, since the inference knobs were swept on both the smaller 010 anchor (0.38333) and the larger i002 checkpoint (0.40590 single-resolution, 0.418265 as the 224+336 ensemble).

Logit adjustment τ and probability threshold T . Table 3 reports the joint sweep over $\tau \in$

$\{0.25, 0.5\}$ and $T \in \{0.03, 0.05, 0.07\}$ on the single-resolution 224-pixel inference recipe (no 336-pixel ensemble). The optimum is sharp: $\tau=0.25$, $T=0.03$ dominates every other cell by between 0.022 and 0.083 Macro F1, and Macro F1 falls off monotonically along both axes. Doubling τ from 0.25 to 0.50 at fixed $T=0.03$ costs 0.022 — evidence that the long-tail boost saturates quickly and that pushing it harder starts emitting too many rare-but-implausible species.

τ	$T=0.03$	$T=0.05$	$T=0.07$
0.25	0.40590	0.36423	0.35069
0.50	0.38440	0.34047	0.32248

Table 3

Joint $\tau \times T$ sweep on the i002 last-four-blocks checkpoint, single-resolution 224-pixel inference. Public Macro F1. The peak at $(\tau=0.25, T=0.03)$ is the anchor value taken into the 224+336 ensemble that produces the headline 0.418265 in Table 5.

Selection rule. Holding $\tau=0.25$ fixed, we compared four selection rules with $k_{\min}=1$, $k_{\max}=10$: *probability threshold* (emit every s with $p_s > T$), *fixed top- k* , *gap* (emit while consecutive species differ by less than a fraction of the top probability), and *relative threshold* (emit every s with p_s within a fraction of the top). Adaptive probability-thresholding at $T=0.03$ wins decisively (Table 4); fixed top- k is second-best at $k=2$, which is consistent with the test set’s empirical mean of ~ 2 –3 species per quadrat; both gap- and relative-threshold rules are well off the pace.

Rule	Value	Public Macro F1
probability threshold (<i>anchor</i>)	$T=0.03$	0.40590
probability threshold	$T=0.05$	0.36423
probability threshold	$T=0.07$	0.35069
fixed top- k	$k=1$	0.28633
fixed top- k	$k=2$	0.39878
fixed top- k	$k=3$	0.39731
fixed top- k	$k=4$	0.38042
fixed top- k	$k=5$	0.35959
gap	0.40	0.33199
gap	0.50	0.33664
gap	0.60	0.34996
relative threshold	0.15	0.31348
relative threshold	0.20	0.31704
relative threshold	0.25	0.31823
relative threshold	0.30	0.33623

Table 4

Selection-rule ablation at $\tau=0.25$ on the i002 last-four-blocks single-resolution checkpoint. Probability-thresholding at $T=0.03$ dominates; fixed top- k is competitive at $k=2$ because the test-set mean cardinality is empirically ~ 2 –3 species per quadrat. Both gap and relative-threshold variants underperform the anchor by between 0.06 and 0.10 Macro F1.

Per-tile aggregation (on the 010 anchor; see ablation notes in `archive/docs/experiments_summary.md`). We tested per-species softmax-mean against a per-species max and a mixture $\alpha \cdot \text{softmax_mean} + (1-\alpha) \cdot \text{noisy_or}$ across $\alpha \in \{0.25, 0.50, 0.75, 1.0\}$. All variants land within 0.0005 Macro F1 of the softmax-mean anchor at 0.38333; the second-best variant, $\alpha=0.75$, scores 0.38278. We read this as evidence that the choice of per-tile aggregation is *not* the bottleneck on this benchmark, and report softmax-mean throughout the rest of the paper without further reservation.

Resolution ensemble. Going from a single-resolution 224-pixel inference at $T=0.05$ to a 224 + 336 probability-averaged ensemble at $\tau=0.25$, $T=0.03$ lifts public Macro F1 from 0.41165 to 0.418265 ($\Delta = +0.0066$) and private Macro F1 from 0.38132 to 0.40283 ($\Delta = +0.022$). The private gain is

more than $3\times$ the public gain — the 336-pixel pass appears to recover rare-species pixels that the single-resolution recipe drops, and those species are the ones the private split weights most heavily.

Takeaway. Among the inference-side knobs, only the logit-adjustment τ and the probability threshold T are load-bearing in the sense that misadjusting either by a single grid step costs more than the run-to-run noise floor of ~ 0.002 Macro F1. The selection rule choice is qualitatively important (probability-threshold beats fixed top- k by ~ 0.007 at the optimum), but inside the probability-threshold family the surface is gentle. The per-tile aggregation is statistically flat at this scale. These observations match the broader pattern of §5: the bottleneck on this benchmark is the visual model’s exposure to multi-species scenes during training, not any of the inference-side hyperparameter choices.

4.8. Aggregate Results

Config	Macro-F1	Δ
BC2.5 last 4 blocks, 224+336+LA, $T=0.03$	0.418265	—
+ cRT triple, $w=0.3$	0.38195	−0.036
+ Genus $\alpha_g=0.1$	0.41246	−0.006
+ Pheno., $\beta=1.0$	0.41346	−0.005

Table 5

All four submitted configurations on the PlantCLEF 2026 hidden test set. The cRT pivot is strongly negative; the genus rerank and the phenology prior move the anchor by -0.006 and -0.005 respectively — magnitudes comfortably inside the run-to-run leaderboard noise floor we previously characterised at ~ 0.002 but consistent in sign. We read this as the BioCLIP 2.5 posterior implicitly absorbing both signals through training-distribution co-occurrence: neither prior carries enough orthogonal information to break the visual anchor.




Figure 9: Public Macro F1 of the four submitted configurations from Table 5. The dashed grey line marks the anchor (0.418265). cRT is the only pivot whose loss exceeds the characterised noise floor; the genus and phenology priors are both within -0.006 of the anchor.

5. Discussion

The pattern in Table 5 reveals an under-appreciated property of strong fine-tuned visual classifiers on biological data: *information-theoretic orthogonality is not the same as statistical orthogonality* once the visual model has been trained on data where the auxiliary signal is naturally correlated with image content. cRT was structurally orthogonal but introduced wrong evidence and dragged the anchor down by 0.036; genus reweighting was statistically dependent on the BioCLIP 2.5 posterior itself and moved the score by -0.006 , close to the noise floor; the seasonal phenology prior — the most genuinely orthogonal source we tested, since date is not in the image at all — moved the score by only -0.005 , again close to the noise floor, because the seasonal structure was already implicit in the leaf and flower phenotypes seen during training. The practical implication is that real headroom on this benchmark requires structurally stronger *visual* models (end-to-end fine-tuning that preserves the prior, multi-crop

test-time augmentation, test-time training, synthetic mosaic training distributions) rather than post hoc reweighting with auxiliary priors. A natural next experiment, blocked here by the lack of test site coordinates, is a climate-envelope variant: per-species 12-D bioclim Gaussians fit from training image lat/lons, scored against WorldClim variables at each test site — this would give a prior conditioned on *location* as well as date, which is the one orthogonal axis that visual features cannot recover from the image alone. We also observed a striking decoupling between validation top-5 accuracy on the held-out 10% split of the single-plant manifest and Macro F1 on the multi-species quadrat leaderboard. Across our twelve strongest runs the two metrics correlate at only $r \approx 0.4$ (Figure 10), and within the unfrozen-blocks ablation of Table 1 they invert: validation top-5 accuracy rises monotonically with unfrozen capacity while leaderboard Macro F1 peaks in the middle and collapses at full fine-tune. We read this as direct evidence that the bottleneck on this benchmark is the single-plant to multi-species distribution shift, not model capacity on the single-plant distribution itself.




Figure 10: Validation top-5 accuracy on the held-out single-plant split versus public Macro F1 on the multi-species quadrat leaderboard, across our twelve strongest runs. The headline anchor (★, i002 last-four-blocks 224+336+LA ensemble) does not sit at the maximum validation point; the i003 capped checkpoints (green squares) and the 015 PC24+iNat mix (teal) reach higher validation accuracy but lower leaderboard score. The dashed line indicates the weak monotone trend ($r \approx 0.4$).

6. Final Submitted System

The configuration that produced our best public score is the BioCLIP 2.5 last-four-blocks anchor from experiment i002, submitted as a single backbone deployed under a resolution-ensemble inference recipe rather than as a multi-checkpoint ensemble. The encoder is BioCLIP 2.5 ViT-H/14 with the last four transformer blocks plus the final layer norm and projection unfrozen and fine-tuned on the enlarged i001 manifest (2.65 million single-plant images across 7,806 species; see §3.3) for ten epochs with the joint species/genus/family cross-entropy loss of §3.2. At inference, each quadrat is resized to 1792×1792 pixels, partitioned into a 4×4 grid of non-overlapping 448-pixel tiles, and each tile is forwarded through the encoder at both 224 and 336 pixels. Per-tile softmax probabilities are aggregated by species-wise mean across the sixteen tiles, the 224 and 336 posteriors are averaged, logit adjustment $-\tau \log \tilde{\pi}_s$ is applied with $\tau=0.25$ (where $\tilde{\pi}_s$ is the Laplace-smoothed training prior for species s), and we emit every species whose post-adjusted probability exceeds $T=0.03$ subject to $k_{\min}=2$ and $k_{\max}=10$. Of the four configurations we submitted, the genus/family rerank and the phenology prior land -0.006 and -0.005 below the anchor respectively (magnitudes inside our characterised noise floor); the triple-backbone cRT mix falls -0.036 below. We accordingly take the unaugmented anchor as our headline submission.

7. Conclusion

We presented the ANU team’s submission to PlantCLEF 2026. Our best public Macro F1 of **0.418265** and private Macro F1 of **0.40283** were obtained from a BioCLIP 2.5 ViT-H/14 fine tune with only the last four transformer blocks (together with the final layer norm and projection) unfrozen, trained for ten epochs on the redownloaded PlantCLEF 2024 manifest (2.65 million single-plant images across 7,806 species, with pre-filled genus and family labels) under a joint species/genus/family cross entropy loss. Inference combined a four-by-four tile decomposition at tile size 448 with no overlap, softmax mean aggregation across the sixteen tiles, an ensemble of inference at 224 and 336 pixel resolutions, logit adjustment ($\tau=0.25$), and an adaptive probability threshold of $T=0.03$. A follow-up experiment (i003) added 161,686 extra images for the 2,022 species with fewer than one hundred training examples and capped every species at five hundred training images; the resulting checkpoint reached 0.40041 on the public leaderboard, below the uncapped anchor by 0.018, which we read as evidence that on this benchmark the marginal value of additional head-class examples is positive and that flattening the head discards information rather than rebalancing it. Three findings carry the bulk of what this benchmark taught us. The first is that partial unfreeze is a sharp, not a plateau, design choice on BioCLIP 2.5: $n=4$ unfrozen blocks is the empirical sweet spot, deviations of one block in either direction cost between 0.009 and 0.014 Macro F1, and a full fine-tune destroys the Tree-of-Life prior and drops the model to 0.208. The second is the validation versus Kaggle decoupling: across our twelve strongest runs the two metrics correlate at only $r \approx 0.4$, and within the unfreeze ablation they invert entirely, which is direct evidence that the bottleneck on this benchmark is the single-plant to multi-species distribution shift rather than model capacity. The third is that auxiliary post-hoc reweighting of an already strong visual posterior carries essentially no usable signal beyond what the visual model has already absorbed. Classifier retraining over a frozen triple-backbone fusion (cRT, $\Delta = -0.036$) introduced wrong evidence and hurt; genus and family co-occurrence priors derived from the same posterior ($\Delta = -0.006$) and a seasonal phenology prior over GBIF observation dates ($\Delta = -0.005$) both moved the anchor by amounts close to our run-to-run noise floor. Fine-tuned visual classifiers on biological data absorb these signals implicitly through training-distribution co-occurrence, leaving little for post-hoc combination to exploit. Several directions for future work follow directly from these findings. The first is to pull more aggressively on the data axis rather than the capacity axis: a synthetic mosaic training regime that composes single-plant images into multi-species scenes with controlled species overlap would let the network see the quadrat distribution during training, which neither the PC24 corpus nor the iNaturalist habitat shots do at scale. The second is self-distillation against our own production checkpoint, using its predictions on the unlabelled LUCAS pseudo-quadrat corpus as soft targets filtered by confidence and entropy. The third is genuinely orthogonal evidence that is not in the image at all, in particular a location-conditioned bioclimatic envelope prior built from training-image coordinates against WorldClim variables at each test site — the one orthogonal axis that visual features cannot recover from the image alone. Finally, per-class threshold calibration on a held-out multi-species validation set would more directly target the Macro F1 metric than the global probability threshold we currently use, and could absorb residual species-frequency imbalances that our manifest cap leaves behind.

Acknowledgments

This work was conducted as part of the COMP3242 Deep Learning course at the Australian National University. We thank the PlantCLEF 2026 organisers for providing the datasets and evaluation infrastructure.

References

- [1] A. Joly, H. Goëau, S. Kahl, L. Pícek, R. R. D. Castaneda, T. Lorieul, B. Deneu, C. Botella, D. Marcos, J. Estopinan, M. Šulc, et al., Overview of LifeCLEF 2024: Evaluation of species distribution

- prediction and identification, in: *Experimental IR Meets Multilinguality, Multimodality, and Interaction*, Lecture Notes in Computer Science, Springer, 2024.
- [2] H. Goëau, A. Joly, P. Bonnet, A. Affouard, B. Deneu, PlantCLEF 2026: Multi-label plant species identification in vegetation plots, Kaggle Competition, 2026. URL: <https://www.kaggle.com/competitions/plantclef-2026>.
 - [3] S. Stevens, J. Wu, M. J. Thompson, E. G. Campolongo, C. H. Song, D. E. Carlyn, L. Dong, W. M. Dahdul, C. Stewart, T. Berger-Wolf, W.-L. Chao, Y. Su, BioCLIP: A vision foundation model for the tree of life, *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* (2024) 19412–19422.
 - [4] H. Goëau, P. Bonnet, A. Joly, The ImageCLEF 2013 plant identification task, in: *CLEF 2013 Working Notes*, 2013.
 - [5] S. H. Lee, C. S. Chan, S. J. Mayo, P. Remagnino, How deep learning extracts and learns leaf features for plant classification, *Pattern Recognition* 71 (2017) 1–13.
 - [6] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, N. Houlsby, An image is worth 16x16 words: Transformers for image recognition at scale, in: *International Conference on Learning Representations (ICLR)*, 2021.
 - [7] H. Goëau, P. Bonnet, A. Joly, Overview of PlantCLEF 2024: Multi-label identification in vegetation plots, in: *CLEF 2024 Working Notes*, *CEUR Workshop Proceedings*, 2024.
 - [8] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, M. Assran, N. Ballas, W. Galuba, R. Howes, P.-Y. Huang, S.-W. Li, I. Misra, M. Rabbat, V. Sharma, G. Synnaeve, H. Xu, H. Jégou, J. Mairal, P. Labatut, A. Joulin, P. Bojanowski, DINOv2: Learning robust visual features without supervision, *Transactions on Machine Learning Research* (2024).
 - [9] S. Woo, S. Debnath, R. Hu, X. Chen, Z. Liu, I. S. Kweon, S. Xie, ConvNeXt V2: Co-designing and scaling convnets with masked autoencoders, in: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023, pp. 16133–16142.
 - [10] E. Ben-Baruch, T. Ridnik, N. Zamir, A. Noy, I. Friedman, M. Protter, L. Zelnik-Manor, Asymmetric loss for multi-label classification, in: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021, pp. 82–91.
 - [11] T.-Y. Lin, P. Goyal, R. Girshick, K. He, P. Dollár, Focal loss for dense object detection, in: *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2017, pp. 2980–2988.
 - [12] A. K. Menon, S. Jayasumana, A. S. Rawat, H. Jain, A. Veit, S. Kumar, Long-tail learning via logit adjustment, in: *International Conference on Learning Representations (ICLR)*, 2021.